

Neural Networks for Natural Language Processing

Jun.-Prof. Sophie Fellenz

Week 4 – Language Modeling and Neural Networks

11 Nov 2024

Questions

- Is Skipgram based on neural networks?
- What is the difference between a neural network word embedding and Skipgram? (question from last exam)

Answers

- Difference:
 - Skipgram embeddings are used inside neural networks (first layer)
 - NN embeddings are layers inside the neural network (usually the last layer(s))
- Therefore:
 - For Skipgram you don't need a neural network (skipgram itself is a log-linear model. The parameters of this model **are** the embeddings)
 - For NN embeddings, you need the network to get the embeddings, the parameters of the NN are **not** the embeddings

Language Models

Are these sentences ok?

- Jane went to the store.
- store to Jane went the.
- Jane went store.
- Jane goed to the store.
- The store went to Jane.
- The food truck went to Jane.

Are these sentences ok?

- Jane went to the store.
- store to Jane went the.
- Jane went store.
- Jane goed to the store.
- The store went to Jane.
- The food truck went to Jane.



Create a grammar of the language

Consider morphology and exceptions

Semantic categories, preferences

And their exceptions

Probabilistic language models

$$P(X) = \prod_{i=1}^I P(\underbrace{x_i}_{\text{Next word}} | \underbrace{x_1, \dots, x_{i-1}}_{\text{Context}})$$

The big problem: How do we predict

$$P(x_i | x_1, \dots, x_{i-1})$$

??

What can we do with LMs?

Score sentences:

- Jane went to the store . -> high
- Store to Jane went the . -> low
- (same as calculating loss for training)

Generate sentences:

while didn't choose end-of-sentence symbol:

calculate probability

sample a new word from the probability distribution

Count-based Language Models

Count-based unigram model

Independence assumption: $P(x_i | x_1, \dots, x_{i-1}) \approx P(x_i)$

Count-based maximum-likelihood estimation:

$$P_{MLE}(x_i) = \frac{c_{train}(x_i)}{\sum_{\tilde{x}} c_{train}(\tilde{x})}$$

Interpolation with UNK model:

$$P(x_i) = (1 - \lambda_{unk}) * P_{MLE}(x_i) + \lambda_{unk} * P_{unk}(x_i)$$

Higher-order n-gram models

Limit context length to n , count, and divide

$$P_{ML}(x_i | x_{i-n+1}, \dots, x_{i-1}) := \frac{c(x_{i-n+1}, \dots, x_i)}{c(x_{i-n+1}, \dots, x_{i-1})}$$
$$P(\text{example} | \text{this is an}) = \frac{c(\text{this is an example})}{c(\text{this is an})}$$

Add smoothing, to deal with zero counts:

$$P(x_i | x_{i-n+1}, \dots, x_{i-1})$$
$$= \lambda P_{ML}(x_i | x_{i-n+1}, \dots, x_{i-1}) + (1 - \lambda) P(x_i | x_{i-n+2}, \dots, x_{i-1})$$

Smoothing methods

Additive/Dirichlet:

Fallback distribution

$$P(x_i | x_{i-n+1}, \dots, x_{i-1}) := \frac{c(x_{i-n+1}, \dots, x_i) + \alpha P(x_i | x_{i-n+2}, \dots, x_{i-1})}{c(x_{i-n+1}, \dots, x_{i-1}) + \alpha}$$

Interpolation hyperparameter

Discounting:

$$P(x_i | x_{i-n+1}, \dots, x_{i-1}) := \frac{c(x_{i-n+1}, \dots, x_i) - d + \alpha P(x_i | x_{i-n+2}, \dots, x_{i-1})}{c(x_{i-n+1}, \dots, x_{i-1})}$$

discount hyperparameter

Interpolation calculated by sum of discounts $\alpha = \sum_{\{\tilde{x}; c(x_{i-n+1}, \dots, \tilde{x}) > 0\}} d$

Problems and solutions?

Cannot share strength among similar words

she bought a car

she purchased a car

she bought a bicycle

she purchased a bicycle

solution: class based language models

Cannot condition on context with intervening words

Dr. Jane Smith

Dr. Gertrude Smith

solution: skip-gram language models

Problems and solutions?

Cannot handle long-distance dependencies

For tennis class he wanted to buy his own racket
for programming class he wanted to buy his own computer

solution: cache, trigger, topic, syntactic models, etc.

When to use n-gram models

- Neural language models (next) achieve better performance, but
- n-gram models are extremely fast to estimate/apply
- n-gram models can be better at modeling low-frequency phenomena

LM Evaluation

Evaluation of LMs

Log-likelihood:

$$LL(\mathcal{E}_{test}) = \sum_{E \in \mathcal{E}_{test}} \log P(E)$$

Per-word Log Likelihood:

$$WLL(\mathcal{E}_{test}) = \frac{1}{\sum_{E \in \mathcal{E}_{test}} |E|} \sum_{E \in \mathcal{E}_{test}} \log P(E)$$

Per-word Entropy:

$$H(\mathcal{E}_{test}) = \frac{1}{\sum_{E \in \mathcal{E}_{test}} |E|} \sum_{E \in \mathcal{E}_{test}} -\log_2 P(E)$$

Perplexity:

$$\text{ppl}(\mathcal{E}_{test}) = 2^{H(\mathcal{E}_{test})} = e^{-WLL(\mathcal{E}_{test})}$$

Evaluation and Vocabulary

- Important: the vocabulary must be the same over models you compare
- Or more accurately, all models must be able to generate the test set
- E.g.: A model that has a vocabulary with only the unknown word is trivial and it would be unfair to compare with a model that has all different words in the vocabulary

Log-linear models


An Alternative: Featurized Models

- Calculate features of the context
- Based on the features, calculate probabilities
- Optimize feature weights using gradient descent, etc.

An Alternative: Featurized Models

Calculate features of the context, calculate probabilities

Each vector is size of
output vocabulary



Feature weights optimized by SGD, etc

Example:

Previous words: “giving a”

a
the
talk
gift
hat
...

$$b = \begin{pmatrix} 3.0 \\ 2.5 \\ -0.2 \\ 0.1 \\ 1.2 \\ \dots \end{pmatrix}$$

Words we're
predicting

How likely
are they?

$$W_{1,a} = \begin{pmatrix} -6.0 \\ -5.1 \\ 0.2 \\ 0.1 \\ 0.5 \\ \dots \end{pmatrix}$$

How likely
are they
given prev.
word is “a”?

$$W_{2,giving} = \begin{pmatrix} -0.2 \\ -0.3 \\ 1.0 \\ 2.0 \\ -1.2 \\ \dots \end{pmatrix}$$

How likely
are they
given 2nd prev.
word is “giving”?

$$s = \begin{pmatrix} -3.2 \\ -2.9 \\ 1.0 \\ 2.2 \\ 0.6 \\ \dots \end{pmatrix}$$

Total score

Training Algorithm

- Calculate the gradient of the loss function with respect to the parameters
- How? Use the chain rule / back-propagation. More in a second
- Update to move in a direction that decreases the loss

What Problems are Handled?

Cannot share strength among similar words

she bought a car
she purchased a car

she bought a bicycle
she purchased a bicycle

not solved yet




Cannot condition on context with intervening words

Dr. Jane Smith

Dr. Gertrude Smith

Solved!



Problems and solutions?

Cannot handle long-distance dependencies

For **tennis** class he wanted to buy his own **raquet**
for **programming** class he wanted to buy his own **computer**

Not solved yet



Beyond linear models

Linear Models can't Learn Feature Combinations

Students take tests → high

Students write tests → low


Teachers take tests → low

Teachers write tests → high


- These can't be expressed by linear features
- What can we do?
 - Remember combinations as features (individual scores for “students take”, “teachers write”)
 - → Feature space explosion!
- Neural networks!

“Neural” Nets


Original Motivation: Neurons in the Brain



Current Conception: Computation Graphs



Input neurons:



Output neurons:


$$x_1 w_{1,1} + x_2 w_{2,1} + x_3 w_{3,1} + b_1$$

$$x_1 w_{1,2} + x_2 w_{2,2} + x_3 w_{3,2} + b_2$$

$$x_1 w_{1,3} + x_2 w_{2,3} + x_3 w_{3,3} + b_3$$

$$x_1 w_{1,4} + x_2 w_{2,4} + x_3 w_{3,4} + b_4$$

Input neurons:



Output neurons:

$$x_1 w_{1,1} + x_2 w_{2,1} + x_3 w_{3,1} + b_1$$

$$x_1 w_{1,2} + x_2 w_{2,2} + x_3 w_{3,2} + b_2$$

$$x_1 w_{1,3} + x_2 w_{2,3} + x_3 w_{3,3} + b_3$$


$$x_1 w_{1,4} + x_2 w_{2,4} + x_3 w_{3,4} + b_4$$

$$\mathbf{z}^T = \mathbf{x}^T \mathbf{W} + \mathbf{b}^T$$

$$\mathbf{z}^T = (z_1, z_2, z_3, z_4), \mathbf{x}^T = (x_1, x_2, x_3),$$

$$\mathbf{W} = (w_{i,j})_{i,j}, \mathbf{b}^T = (b_1, b_2, b_3, b_4)$$

Input neurons:



Output neurons:

$$x_1 w_{1,1} + x_2 w_{2,1} + x_3 w_{3,1} + b_1$$

$$x_1 w_{1,2} + x_2 w_{2,2} + x_3 w_{3,2} + b_2$$

$$x_1 w_{1,3} + x_2 w_{2,3} + x_3 w_{3,3} + b_3$$

$$x_1 w_{1,4} + x_2 w_{2,4} + x_3 w_{3,4} + b_4$$

Neurons in the brain only activate after a certain threshold is overcome

Simulate with activation function $\sigma(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$

$$\text{Output} = \sigma(x^T W + b^T)$$

Input layer

Hidden layers

Output




Usually we use multiple layers

$$\textit{Output} = \sigma(\dots \sigma(x^T W_1 + b_1^T) \dots W_n + b_n^T)$$

Input layer

Hidden layers

Output



How does this work specifically?

Example: Given cat or dog image.

Task: Find out if it is a cat or a dog. Output (1,0) if cat and (0,1) if dog.

Input layer

Hidden layers

Output



Why does this work?

- One can prove that any function, i.e. the cat-dog recognition function can be approximated by a neural network!
- Network needs to be large enough and have correct weights W and b

Input layer

Hidden layers

Output



How do we find out how large it needs to be?

- Experimentation!

How do we find the weights?

- We initiate randomly and then train on train data!

Training


- Given:
Training data, i.e. input data x where desired output y is known.
Neural Network $n_{W,b}$
- Output of Neural Network: $n_{W,b}(x)$
Desired output: y
- Loss: $\|n_{W,b}(x) - y\|^2$
- Goal: Minimize loss by optimizing weights W and b

expression:


 x

graph:

A **node** is a {tensor, matrix, vector, scalar} value



- An **edge** represents a function argument (and also a data dependency). They are just pointers to nodes.
- A **node** with an incoming **edge** is a **function** of that edge's tail node.
- A **node** knows how to compute its value and the value of its derivative w.r.t each argument (edge) times a derivative of an arbitrary input $\frac{\partial f}{\partial u}$.




expression:

$$x^T W + b$$

graph:

Functions can be nullary, unary, binary, ... n-ary.


Often they are unary or binary



expression:

$$\sigma(x^T W + b^T)$$


graph:



expression:

$$\sigma(x^T W + b^T) - y^T$$


graph:



expression:

$$\text{Loss} = \|\sigma(x^T W + b^T) - y^T\|^2 = f_5(f_4(f_3(f_2(f_1(x, W), b)), y))$$

graph:




Algorithms (1)


- Graph construction
- Forward propagation

In topological order, compute the **value** of the node given its inputs


Forward Propagation



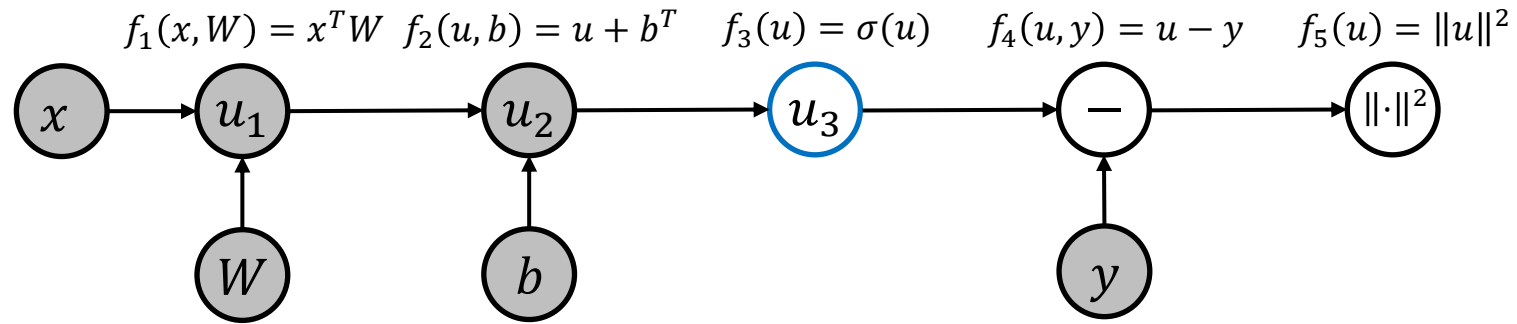
Forward Propagation




Forward Propagation




Forward Propagation



Forward Propagation



Forward Propagation



Algorithms (2)

- Aim: Minimize loss $f(x, W, b) = \|\sigma(x^T W + b^T) - y^T\|^2$
w.r.t. weights W, b
- Idea: Gradient = Direction of highest increase
Calculate its gradients $\frac{\partial f}{\partial W}$ and $\frac{\partial f}{\partial b}$ and move against it
- Parameter update:
 - Move the parameters against the direction of this derivative
 - $W = W - \alpha * \frac{\partial f}{\partial W}, b = b - \alpha * \frac{\partial f}{\partial b}$
- $\alpha > 0$ is learning rate

Algorithms (2)

Back-propagation:


- Process examples in reverse topological order
- Calculate the derivatives of the parameters with respect to the final value

Back Propagation

Expression:

$$\text{Loss} = \|\sigma(x^T W + b^T) - y^T\|^2$$


Aim: Minimize loss by optimizing weights W, b



Back Propagation


Step 1: Compute all derivatives at every node wrt. the inputs
(at relevant edges)

Example: $\frac{\partial f_2}{\partial b} = 1$



Back Propagation

Step 1: Compute all derivatives at every node wrt. the inputs.



Back Propagation


Step 2: Use the chain rule.

Example: $f(x, W, b) = \|\sigma(x^T W + b^T) - y^T\|^2$

$= f_5(f_4(f_3(f_2(f_1(x, W), b)), y))$

$$\frac{\partial f(x, W, b)}{\partial b} = \frac{\partial f_5}{\partial u_4}(u_4) \frac{\partial f_4}{\partial b}(u_3, y) = 2u_4 \frac{\partial f_4}{\partial b}(u_3, y)$$

$$f_1(x, W) = x^T W \quad f_2(u, b) = u + b^T \quad f_3(u) = \sigma(u) \quad f_4(u, y) = u - y \quad f_5(u) = \|u\|^2$$



Back Propagation


Step 2: Use the chain rule.

Example: $f(x, W, b) = \|\sigma(x^T W + b^T) - y^T\|^2$

$$= f_5(f_4(f_3(f_2(f_1(x, W), b)), y))$$

$$\frac{\partial f(x, W, b)}{\partial b} = \frac{\partial f_5}{\partial u_4}(u_4) \frac{\partial f_4}{\partial u_3}(u_3, y) \frac{\partial f_3}{\partial b}(u_2) = 2u_4 \frac{\partial f_3}{\partial b}(u_2)$$

$$f_1(x, W) = x^T W \quad f_2(u, b) = u + b^T \quad f_3(u) = \sigma(u) \quad f_4(u, y) = u - y \quad f_5(u) = \|u\|^2$$




Back Propagation

Step 2: Use the chain rule.

Example: $f(x, W, b) = \|\sigma(x^T W + b^T) - y^T\|^2 = f_5(f_4(f_3(f_2(f_1(x, W), b)), y))$

$$\frac{\partial f(x, W, b)}{\partial b} = \frac{\partial f_5}{\partial u_4}(u_4) \frac{\partial f_4}{\partial u_3}(u_3, y) \frac{\partial f_3}{\partial u_2}(u_2) \frac{\partial f_2}{\partial b}(u_1, b) = 2u_4 \sigma'(u_2)$$




Back Propagation

Step 2: Use the chain rule.

Similarly, we calculate $\frac{\partial f}{\partial W}$. Note that we already have calculated $\frac{\partial f}{\partial u_2}$ previously.

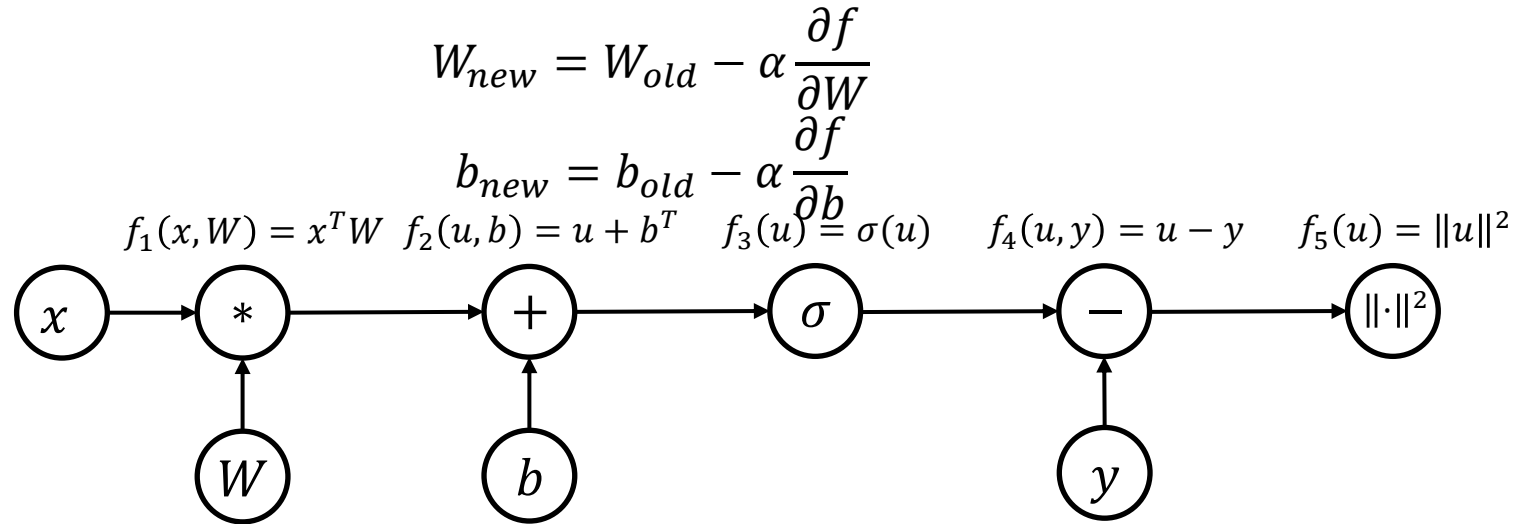
In summary: Multiply derivatives starting from end



Back Propagation

Step 3: Apply Gradient descent


Update: $\alpha > 0$ learning rate



Back to language modeling

Feed-forward Neural Language Models

- See Bengio et al. 2003



Example of Combination Features

- Word embeddings capture features of words
- e.g. feature 1 indicates verbs, feature 2 indicates determiners
- A row in the weight matrix (together with the bias) can capture particular combinations of these features
- e.g. the 34th row in the weight matrix looks at feature 1 in the second to-previous word, and feature 2 in the previous word

giving

a


$$\begin{bmatrix} 1.2 \\ -0.1 \\ 0.7 \\ -2.1 \\ 0.5 \end{bmatrix} * \begin{bmatrix} 1.5 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} + \begin{bmatrix} -0.3 \\ 2.0 \\ 0.6 \\ -0.8 \\ -0.4 \end{bmatrix}$$

w_{34} b_{34}

+ -2

positive number if
the previous word is a
determiner and
second-to-previous
word is a verb

Where is Strength Shared?



What Problems are Handled?

Cannot share strength among similar words

she bought a car
she purchased a car

she bought a bicycle
she purchased a bicycle


solved, and similar contexts as well!

Cannot condition on context with intervening words

Dr. Jane Smith

Dr. Gertrude Smith

Solved!



Problems and solutions?

Cannot handle long-distance dependencies

For **tennis** class he wanted to buy his own **raquet**
for **programming** class he wanted to buy his own **computer**

Not solved yet



Many Other Potential Designs!

- Neural networks allow design of arbitrarily complex functions!
- In future classes:
 - Recurrent neural network LMs
 - Transformer LMs

Next lecture
Recurrent Neural Networks

Acknowledgements

- CMU Advanced NLP Course:
- <https://phontron.com/class/anlp2022/schedule.html>
- Sören Laue
- Feibai Huang

References

- Video on Backprop by Andrej Karpathy:
- <https://youtu.be/VMj-3S1tku0>
- Video on Language modeling by Andrej Karpathy:
- <https://youtu.be/PaCmpygFfXo>